

MPP coursework

Details and answers to common questions

Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.


http://creativecommons.org/licenses/by-nc-sa/4.0/deed.en_US

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Acknowledge EPCC as follows: “© EPCC, The University of Edinburgh, www.epcc.ed.ac.uk”

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

Percolation



Algorithm

	1		2	3
4	5	6	7	
	8		9	
10		11		12
		13		

	5		7	7
5	8	8	9	
	8		9	
10		13		12
		13		

	8		9	9
8	8	9	9	
	8		9	
10		13		12
		13		

	8		9	9
8	9	9	9	
	9		9	
10		13		12
		13		

	9		9	9
9	9	9	9	
	9		9	
10		13		12
		13		

Halos

0	0	0	0	0	0	0
0	0	1	0	1	1	0
0	1	1	1	1	0	0
0	0	1	0	1	0	0
0	1	0	1	0	1	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	1	0	2	3	0
0	4	5	6	7	0	0
0	0	8	0	9	0	0
0	10	0	11	0	12	0
0	0	0	13	0	0	0
0	0	0	0	0	0	0


0	0	0	0	0	0	0
0	0	5	0	7	7	0
0	5	8	8	9	0	0
0	0	8	0	9	0	0
0	10	0	13	0	12	0
0	0	0	13	0	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	8	0	9	9	0
0	8	8	9	9	0	0
0	0	8	0	9	0	0
0	10	0	13	0	12	0
0	0	0	13	0	0	0
0	0	0	0	0	0	0


0	0	0	0	0	0	0
0	0	8	0	9	9	0
0	8	9	9	9	0	0
0	0	9	0	9	0	0
0	10	0	13	0	12	0
0	0	0	13	0	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	9	0	9	9	0
0	9	9	9	9	0	0
0	0	9	0	9	0	0
0	10	0	13	0	12	0
0	0	0	13	0	0	0
0	0	0	0	0	0	0

$L = 432$, $\rho = 0.411$, *seed* = 6543



Largest cluster (just
fails to percolate)



Largest three
clusters

Case study: serial code



- Uses explicit halos for boundary conditions
 - distinguish between L , M and N although all the same values in serial
 - declares arrays to have bounds $0, 1, \dots, M+1$ and $0, 1, \dots, N+1$
 - fills boundaries (all pixels with $i=0$, $i=M+1$, $j=0$ or $j=N+1$) with 0
 - initializes random number generator with seed (a runtime argument)
 - set rock to zero with appropriate density ρ
 - fills empty spaces with unique positive values
 - applies the update equation for many steps
 - copies *old* back to *map* excluding the halos
 - writes *map* to an output file “map.pgm”
- Simplifications
 - map dimension L , density ρ and file name specified at compile time

Case study: parallel algorithm

- Split map into slices amongst all processes
 - divide across the first index i in C, second index j in Fortran
 - local array sizes are now M and N (plus halos)
 - C: $M = L/P$, $N = L$; Fortran: $M = L$, $N = L/P$
- Rank 0 initialises $L \times L$ *map* array (with no halos)
 - scatter to local *smallmap* array (also with no halos)
- Implement exactly the same calculation as before except:
 - **before start** of each step, swap boundary values as appropriate
- Gather *smallmap* array back to *map*
 - Rank 0 writes to file

Domain decomposition

- Different choices in C and Fortran



Coursework

- Differences from Case Study
 - use a 2D decomposition, i.e. try to split into many rectangular regions, not restricted to slices
 - requires more halo swaps than before (and they must be non-blocking!)
 - you cannot do the distribution and collection using scatter and gather
 - different boundary conditions
 - periodic in *second* dimension: $map_{i,0} = map_{i,N}$; $map_{i,N+1} = map_{i,1}$
 - compute the average value of map at regular intervals
 - stop when $nchange = 0$

Timeline

- Start working on Case Study example as soon as you have completed the other exercises
 - I will hand out simple solutions to the exercises next week
- First coursework due 16:00 GMT on Fri 6th Nov
 - *plan* for how your final 2D coursework solution will be designed, implemented and tested (to improve on 1D Case Study)
 - performance evaluation of 1D Case Study
 - sample 1D solution will be provided
- Second coursework due 16:00 GMT on Wed 2nd Dec
 - source code for your 2D implementation
 - report on the *actual* implementation, testing and performance

Questions

- What if the map does not decompose exactly?
 - you might want to deal with this case (see tutorial problems)
 - however, perfectly acceptable to quit and say “can’t do this!”
 - can still run on many different process counts
- What about a prime number of processes?
 - on 3 processes, for example, your code should just use a 1x3 or 3x1 decomposition (i.e. slices as for the case study)
 - 1D decomp is a subset of 2D: your code should run without problems
- Should I get the same answer?
 - Yes! The exercise has been designed so that (rather unusually for a parallel program) I would expect you to get exactly the same answer regardless of how many processes you run on

Hints (i)

- Do no worry about doing elegant IO
 - just implement something simple that works
 - e.g. broadcast the whole map to all processes
 - copy relevant section to your local smallmap array
- Test section by section
 - run on a single process
 - checks for serious bugs!
 - run for zero steps
 - checks the initialisation, distribution, collection and IO are correct
 - run with no halo swaps
 - checks the calculation is correct
 - run with only horizontal or vertical halo swaps
 - ...

Hints (ii)

- You don't always need to run for thousands of steps
 - can set a maximum count for performance testing
- You are welcome to experiment with larger or smaller maps
- Bring your problems to the tutorial and drop-in sessions!